

Make Sense of Deep Neural Networks using TensorBoard

github.com/PythonWorkshop/tensorboard_demos

Arpan Chakraborty

[@runoffthemill](#)
github.com/napratin



UDACITY



Getting Started

Note: This notebook is written in Python 3. You'll need Jupyter/iPython to run it.

Local installation

Fetch the repo: github.com/PythonWorkshop/tensorboard_demos

```
git clone git@github.com:PythonWorkshop/tensorboard_demos
cd tensorboard_demos/
```

Option A: Conda install

```
conda env create
source activate tensorflow
jupyter notebook tensorboard_basics.ipynb
```

Option B: Pip install

```
pip3 install numpy matplotlib scikit-learn
pip3 install --upgrade <binary URL for your system>
jupyter notebook tensorboard_basics.ipynb
```

See [TensorFlow instructions](#) to pick the correct binary URL for your system.



Binder

launch binder

If you have trouble getting TensorFlow to work, hit the **launch binder** badge to run in the cloud. Note that this is an experimental feature.



Neural Networks Primer

A brief introduction to neural networks, with an eye on data structures.



Simple linear model

Express output (y) as a linear combination of inputs (x):

$$y = x_0 w_0 + x_1 w_1 + \dots + x_n w_n + b$$

Learning goal: Given enough examples of x and y , find w and b that best capture the mapping.



Simple linear model

Express output (y) as a linear combination of inputs (x):

$$y = x_0 w_0 + x_1 w_1 + \dots + x_n w_n + b$$


Learning goal: Given enough examples of x and y , find w and b that best capture the mapping.

Linear model: Vector notation


$$y = \begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_n \end{bmatrix} + b$$



Linear model: Graphical representation



Single neuron model



Neuron model: Vector notation

$$y = f\left(\begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_n \end{bmatrix} + b\right)$$

Here $f()$ is a *non-linear* activation function, that maps real-valued output y to some target space, e.g. the sigmoid function maps to $[0, 1]$.



Multiple neurons

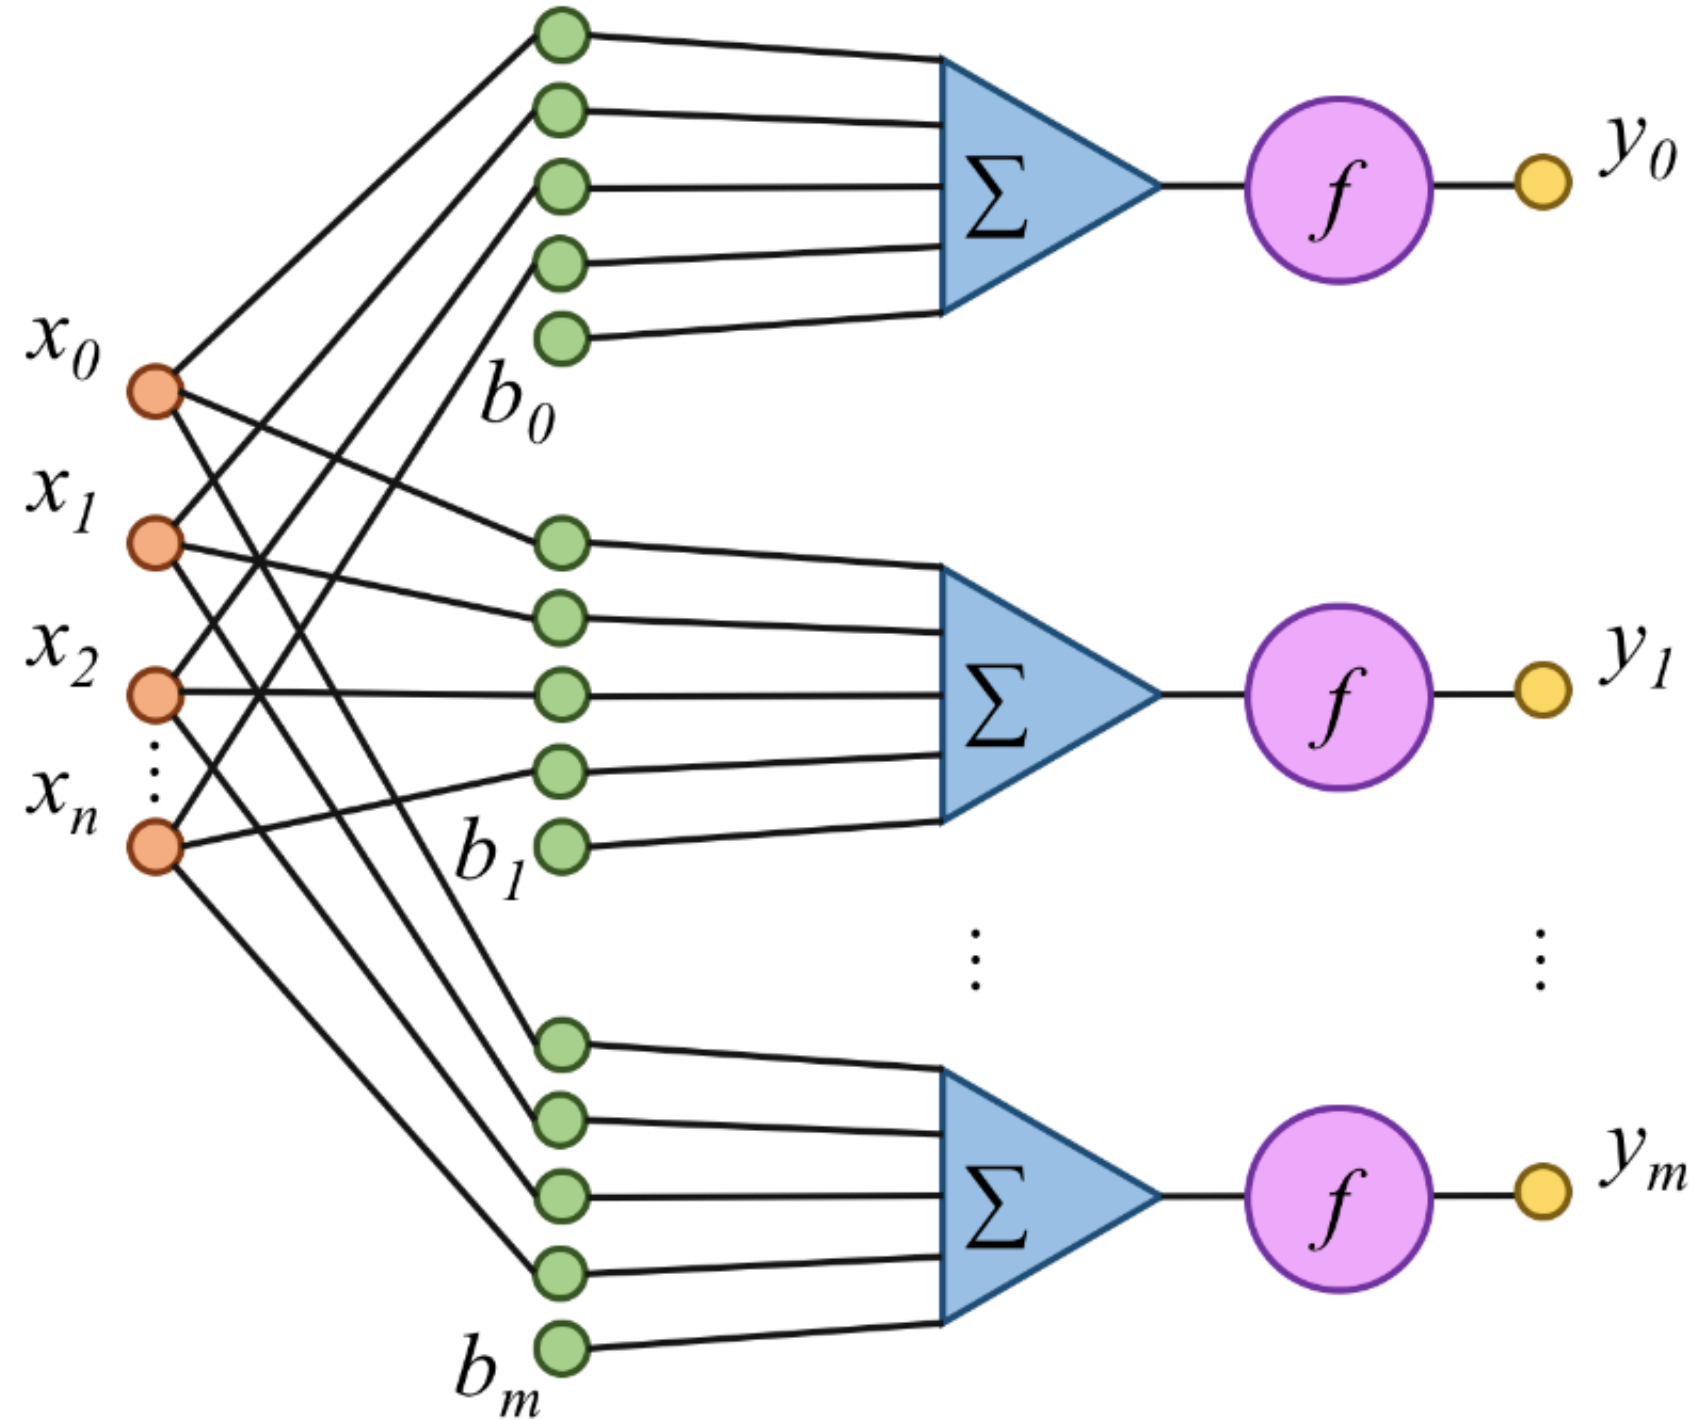
Instead of a single neuron, now we have m of them.

$$\begin{bmatrix} y_0 & y_1 & \dots & y_m \end{bmatrix} = f\left(\begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0m} \\ w_{10} & w_{11} & \dots & w_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n0} & w_{n1} & \dots & w_{nm} \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \dots & b_m \end{bmatrix}\right)$$

This is equivalent to a *layer* of neurons.



Multiple neurons: Graphical representation



Process multiple data samples at once

Instead of a single data sample, now we are dealing with k samples in parallel.

$$\begin{bmatrix} y_{00} & y_{01} & \dots & y_{0m} \\ y_{10} & y_{11} & \dots & y_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ y_{k0} & y_{k1} & \dots & y_{km} \end{bmatrix} = f\left(\begin{bmatrix} x_{00} & x_{01} & \dots & x_{0n} \\ x_{10} & x_{11} & \dots & x_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{k0} & x_{k1} & \dots & x_{kn} \end{bmatrix} \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0m} \\ w_{10} & w_{11} & \dots & w_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n0} & w_{n1} & \dots & w_{nm} \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \dots & b_m \end{bmatrix} \right)$$

Note: Here the bias values b are repeated for each sample (row).



Process multiple data samples at once

Instead of a single data sample, now we are dealing with k samples in parallel.

$$\begin{bmatrix} y_{00} & y_{01} & \dots & y_{0m} \\ y_{10} & y_{11} & \dots & y_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ y_{k0} & y_{k1} & \dots & y_{km} \end{bmatrix} = f\left(\begin{bmatrix} x_{00} & x_{01} & \dots & x_{0n} \\ x_{10} & x_{11} & \dots & x_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{k0} & x_{k1} & \dots & x_{kn} \end{bmatrix} \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0m} \\ w_{10} & w_{11} & \dots & w_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n0} & w_{n1} & \dots & w_{nm} \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \dots & b_m \end{bmatrix} \right)$$

Note: Here the bias values b are repeated for each sample (row).

What does this all look like?



Process multiple data samples at once

Instead of a single data sample, now we are dealing with k samples in parallel.

$$\begin{bmatrix} y_{00} & y_{01} & \dots & y_{0m} \\ y_{10} & y_{11} & \dots & y_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ y_{k0} & y_{k1} & \dots & y_{km} \end{bmatrix} = f\left(\begin{bmatrix} x_{00} & x_{01} & \dots & x_{0n} \\ x_{10} & x_{11} & \dots & x_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{k0} & x_{k1} & \dots & x_{kn} \end{bmatrix} \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0m} \\ w_{10} & w_{11} & \dots & w_{1m} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n0} & w_{n1} & \dots & w_{nm} \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \dots & b_m \end{bmatrix} \right)$$


Note: Here the bias values b are repeated for each sample (row).


What does this all look like?

TENSORS!



Deep Neural Networks





Ready for some Machine Learning?



Generate Some Data to Classify

```
In [1]: import numpy as np

num_examples, num_features = (1000, 2) # dataset size
num_classes = 2 # binary classification task

X = np.random.random((num_examples, num_features))
y = np.int_(X[:, 0] * X[:, 0] + X[:, 1] >= 1).reshape(-1, 1)
```



```
In [2]: print("Features (X): {shape[0]}x{shape[1]}".format(shape=X.shape))  
print(X[:10])
```

```
Features (X): 1000x2  
[[ 0.58578979  0.96048674]  
 [ 0.41171871  0.52405548]  
 [ 0.50849822  0.75811416]  
 [ 0.6386604   0.23468108]  
 [ 0.14959901  0.54418769]  
 [ 0.29773577  0.72048556]  
 [ 0.67330963  0.31292161]  
 [ 0.87784032  0.74873901]  
 [ 0.60013032  0.35332369]  
 [ 0.9033483   0.74016079]]
```



```
In [3]: print("Labels (y): {shape[0]}".format(shape=y.shape))  
        print(y[:10])
```

```
Labels (y): 1000
```

```
[[1]  
 [0]  
 [1]  
 [0]  
 [0]  
 [0]  
 [0]  
 [0]  
 [1]  
 [0]  
 [1]]
```



Visualize Data


```
In [4]: import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')

def plot_data_2D(X, **kwargs):
    """Plot 2D data points in X = np.array([(x0, y0), (x1, y1), ...])."""
    fig, ax = plt.subplots(figsize=(9, 9))
    ax.scatter(X[:, 0], X[:, 1],
               s=35, cmap=plt.cm.get_cmap('rainbow', num_classes), **kwargs)
    ax.set_aspect('equal')
    ax.set_xlim(0, 1)
    ax.set_ylim(0, 1)
    return fig, ax
```



```
In [5]: plot_data_2D(X, c=y)
```

```
Out[5]: (<matplotlib.figure.Figure at 0x10986c048>,  
<matplotlib.axes._subplots.AxesSubplot at 0x109863b00>)
```



Prepare Data for Training and Testing

```
In [6]: # Split data into training and test sets
        from sklearn.cross_validation import train_test_split


        test_size = 300
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)
        print("Split dataset: {} training, {} test samples".format(len(X_train),
                                                                    len(X_test)))
```

Split dataset: 700 training, 300 test samples




```
In [7]: # Plot training data  
plot_data_2D(X_train, c=y_train)
```

```
Out[7]: (<matplotlib.figure.Figure at 0x10afc7f28>,  
<matplotlib.axes._subplots.AxesSubplot at 0x10f2a8320>)
```



```
In [8]: # Plot test data  
plot_data_2D(X_test, c=y_test)
```

```
Out[8]: (<matplotlib.figure.Figure at 0x10f2f7278>,  
<matplotlib.axes._subplots.AxesSubplot at 0x10f30ccf8>)
```



Build Computation Graph

This graph defines the structure of your Neural Network as well as information flow.

```
In [9]: import tensorflow as tf

# Create a TensorFlow session
session = tf.Session()

# Placeholders for input features and ground truth labels
X_placeholder = tf.placeholder(tf.float32,
                              shape=(None, num_features),
                              name="X")
y_placeholder = tf.placeholder(tf.int64,
                              shape=(None, 1),
                              name="y")
```



Define layer creation functions

```
In [11]: def linear_layer(input_tensor, num_units):  
    """Linear activation layer: output = inputs * weights + biases"""  
    num_inputs = input_tensor.get_shape()[1].value # inspect tensor size  
    weights = tf.Variable(  
        tf.truncated_normal([num_inputs, num_units], stddev=1.0),  
        name='weights')  
    biases = tf.Variable(  
        tf.zeros([num_units]),  
        name='biases')  
    return tf.add(tf.matmul(input_tensor, weights, name='multiply'), biases,  
                  name='add')
```



```
In [12]: plot_activation(lambda x: 0.5 + 0.1*x, title='Linear') # example
```

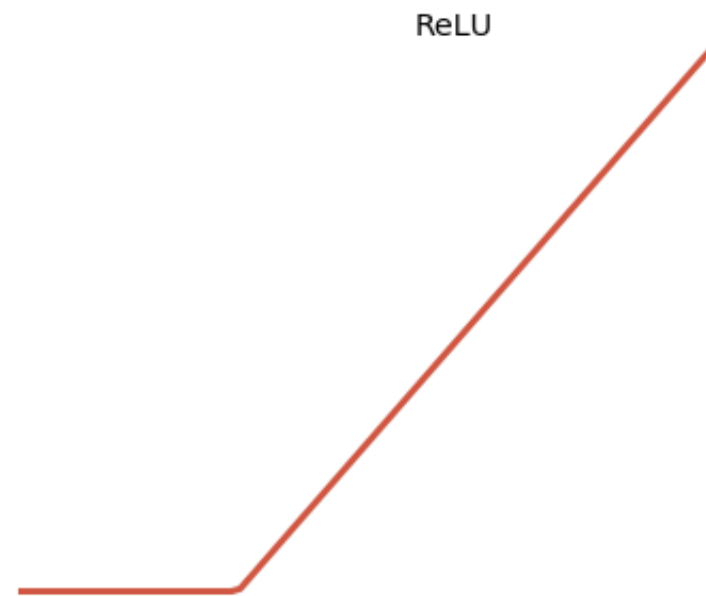


```
In [13]: def sigmoid_layer(input_tensor, num_units):  
        """Sigmoid activation layer: output = sigmoid(inputs * weights + biases)"""  
        return tf.nn.sigmoid(linear_layer(input_tensor, num_units), name='sigmoid')  
  
plot_activation(lambda x: 1.0 / (1.0 + np.exp(-x)), title='Sigmoid') # example
```

Sigmoid



```
In [14]: def relu_layer(input_tensor, num_units):  
         """ReLU activation layer: output = ReLU(inputs * weights + biases)"""  
         return tf.nn.relu(linear_layer(input_tensor, num_units), name='relu')  
  
plot_activation(lambda x: (0.5 + 0.1*x).clip(min=0), title='ReLU') # example
```



Create network structure

```
In [15]: # Let's make a network with 2 hidden layers (and an output layer)
hidden1_num_units = 10
hidden2_num_units = 3
output_num_units = 1 # binary classification needs only 1 output

with tf.name_scope('hidden1'):
    hidden1 = relu_layer(X_placeholder, hidden1_num_units)

with tf.name_scope('hidden2'):
    hidden2 = relu_layer(hidden1, hidden2_num_units)

with tf.name_scope('output'):
    output = sigmoid_layer(hidden2, output_num_units)
```



Choose an error metric

```
In [16]: def l2_loss(logits, labels):  
         """Euclidean distance or L2-norm: sqrt(sum((logits - labels)^2))"""  
         labels = tf.to_float(labels)  
         return tf.nn.l2_loss(logits - labels, name='l2_loss')  
  
         with tf.name_scope('error'):  
             error = l2_loss(output, y_placeholder) # predicted vs. true labels  
             tf.scalar_summary(error.op.name, error) # write error (loss) to log
```



```
In [18]: print("Nodes in computation graph:")
+ "\n Input features: {}".format(X_placeholder)
+ "\n Hidden layer 1: {}".format(hidden1)
+ "\n Hidden layer 2: {}".format(hidden2)
+ "\n Output labels : {}".format(output)
+ "\n Ground truth  : {}".format(y_placeholder)
+ "\n Error metric   : {}".format(error))
```

```
Nodes in computation graph:
Input features: Tensor("X:0", shape=(?, 2), dtype=float32)
Hidden layer 1: Tensor("hidden1/relu:0", shape=(?, 10), dtype=float32)
Hidden layer 2: Tensor("hidden2/relu:0", shape=(?, 3), dtype=float32)
Output labels : Tensor("output/sigmoid:0", shape=(?, 1), dtype=float32)
Ground truth  : Tensor("y:0", shape=(?, 1), dtype=int64)
Error metric   : Tensor("error/l2_loss:0", shape=(), dtype=float32)
```



Setup Logging

These logs are later read and visualized by TensorBoard.

```
In [19]: import time
import os

log_basedir = "logs"
run_label = time.strftime('%Y-%m-%d_%H-%M-%S') # e.g. 2016-08-18_21-30-45
log_path = os.path.join(log_basedir, run_label)
all_summaries = tf.merge_all_summaries()
summary_writer = tf.train.SummaryWriter(log_path, session.graph)
print("Logging to: {}".format(log_path))
```

Logging to: logs/2016-08-20_17-35-19



Train your Model

```
In [20]: # Pick a training algorithm
def sgd_train(error, learning_rate=0.01):
    """Gradient descent optimizer for training.

    Creates an optimizer to compute and apply gradients to all trainable variables.

    Args:
        error: Error (loss) metric.
        learning_rate: Controls the size of each step the optimizer takes.
    Returns:
        training: Training operation, ready to be called with tf.Session.run().
    """
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)
    return optimizer.minimize(error)

with tf.name_scope('training'):
    training = sgd_train(error)
```



```
In [21]: # Define training parameters
num_steps = 1000 # how many iterations to train for
batch_size = 100 # how many samples in each iteration

# Initialize variables
init_op = tf.initialize_all_variables()
session.run(init_op)
```



```
In [22]: # Run training operation for num_steps iterations
for step in range(num_steps):
    # Randomly pick batch_size samples from training set
    sample_idx = np.random.choice(len(X_train), batch_size, replace=False)
    feed_dict = {
        X_placeholder: X_train[sample_idx, :],
        y_placeholder: y_train[sample_idx, :]
    }
    # Note: feed_dict uses placeholder objects as key!

    # Train for one iteration, time it
    start_time = time.time()
    _, error_value = session.run([training, error], feed_dict=feed_dict)
    duration = time.time() - start_time

    # Print an overview and write summaries (logs) every 100 iterations
    if step % 100 == 0 or step == (num_steps - 1):
        print("Step {:4d}: training error = {:.2f} ({:0.3f} sec)".format(
            step, error_value, duration))
        summary_str = session.run(all_summaries, feed_dict=feed_dict)
        summary_writer.add_summary(summary_str, step)
        summary_writer.flush()
```

```
Step    0: training error = 21.90 (0.016 sec)
Step   100: training error =  2.40 (0.004 sec)
Step   200: training error =  3.04 (0.003 sec)
Step   300: training error =  0.97 (0.003 sec)
Step   400: training error =  1.54 (0.002 sec)
Step   500: training error =  0.73 (0.004 sec)
Step   600: training error =  0.56 (0.001 sec)
Step   700: training error =  0.69 (0.002 sec)
Step   800: training error =  2.18 (0.003 sec)
Step   900: training error =  1.28 (0.003 sec)
Step   999: training error =  1.33 (0.003 sec)
```



Test your Model

```
In [23]: # Check performance on test set
y_test_pred, test_error = session.run([output, error], feed_dict={
    X_placeholder: X_test,
    y_placeholder: y_test
})
# Note: The placeholder shapes must be compatible with the tensors being supplied!
```




```
In [24]: y_test_pred = np.int_(np.round_(y_test_pred))
mismatches = (y_test_pred - y_test).flat != 0
print("Test error = {:.2f} ({} mismatches)"
      .format(test_error, sum(np.int_(mismatches))))

_, ax = plot_data_2D(X_test, c=y_test_pred)
ax.scatter(X_test[mismatches, 0], X_test[mismatches, 1],
          s=128, marker='o', facecolors='none', edgecolors='black', linewidths=1.5)
```

Test error = 2.59 (7 mismatches)

Out[24]: <matplotlib.collections.PathCollection at 0x1149c2710>



Visualize using TensorBoard


Run TensorBoard with the following command, passing in the appropriate log directory:

```
tensorboard --logdir=logs
```

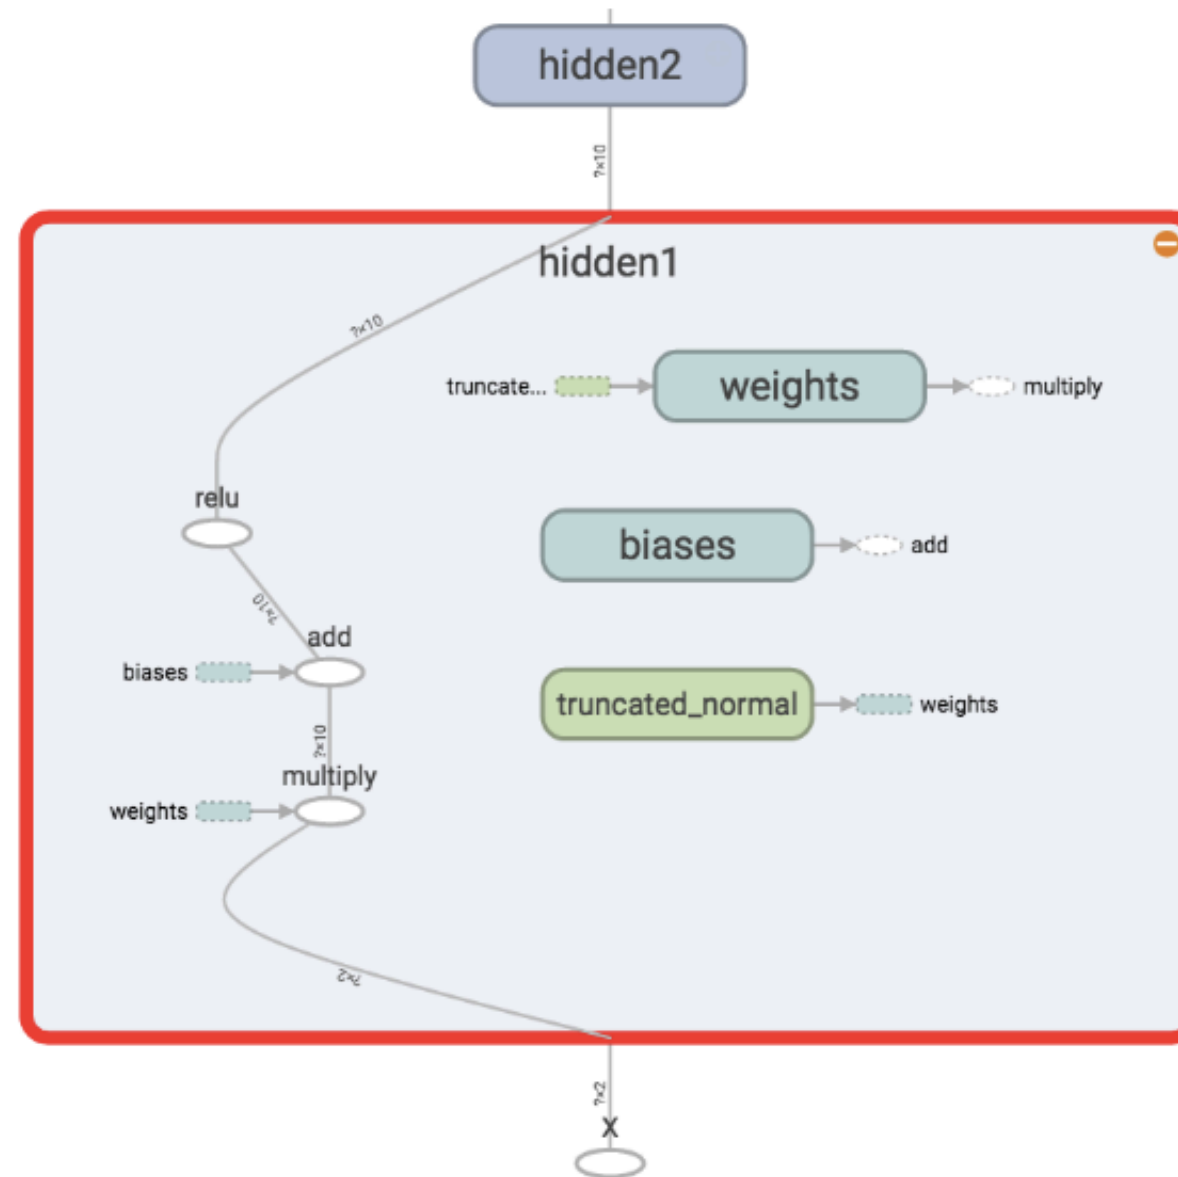
And then open the URL that gets printed, in your browser (typically: <http://0.0.0.0:6006>).



Computation graph



Layer detail: **hidden1**



hidden1 ^

Subgraph: 12 nodes

Attributes (0)

Inputs (1)

○ X 7×2


Outputs (1)

○ hidden2/multiply 7×10

Remove from main graph



Error plot



Summary

- Define a neural network model as nodes of a computation graph in TensorFlow.
- Train your model by presenting labeled data and applying an optimization algorithm.
- Log important values like error during training.
- Test your resulting model on unseen data.
- Visualize computation graph and logged data using TensorBoard.

Note: You can convert this notebook into slides (HTML + reveal.js) and serve them using:

```
jupyter nbconvert tensorboard_basics.ipynb --to slides --post serve
```



Make Sense of Deep Neural Networks using TensorBoard

github.com/PythonWorkshop/tensorboard_demos

Arpan Chakraborty

[@runoffthemill](#)
github.com/napratin

